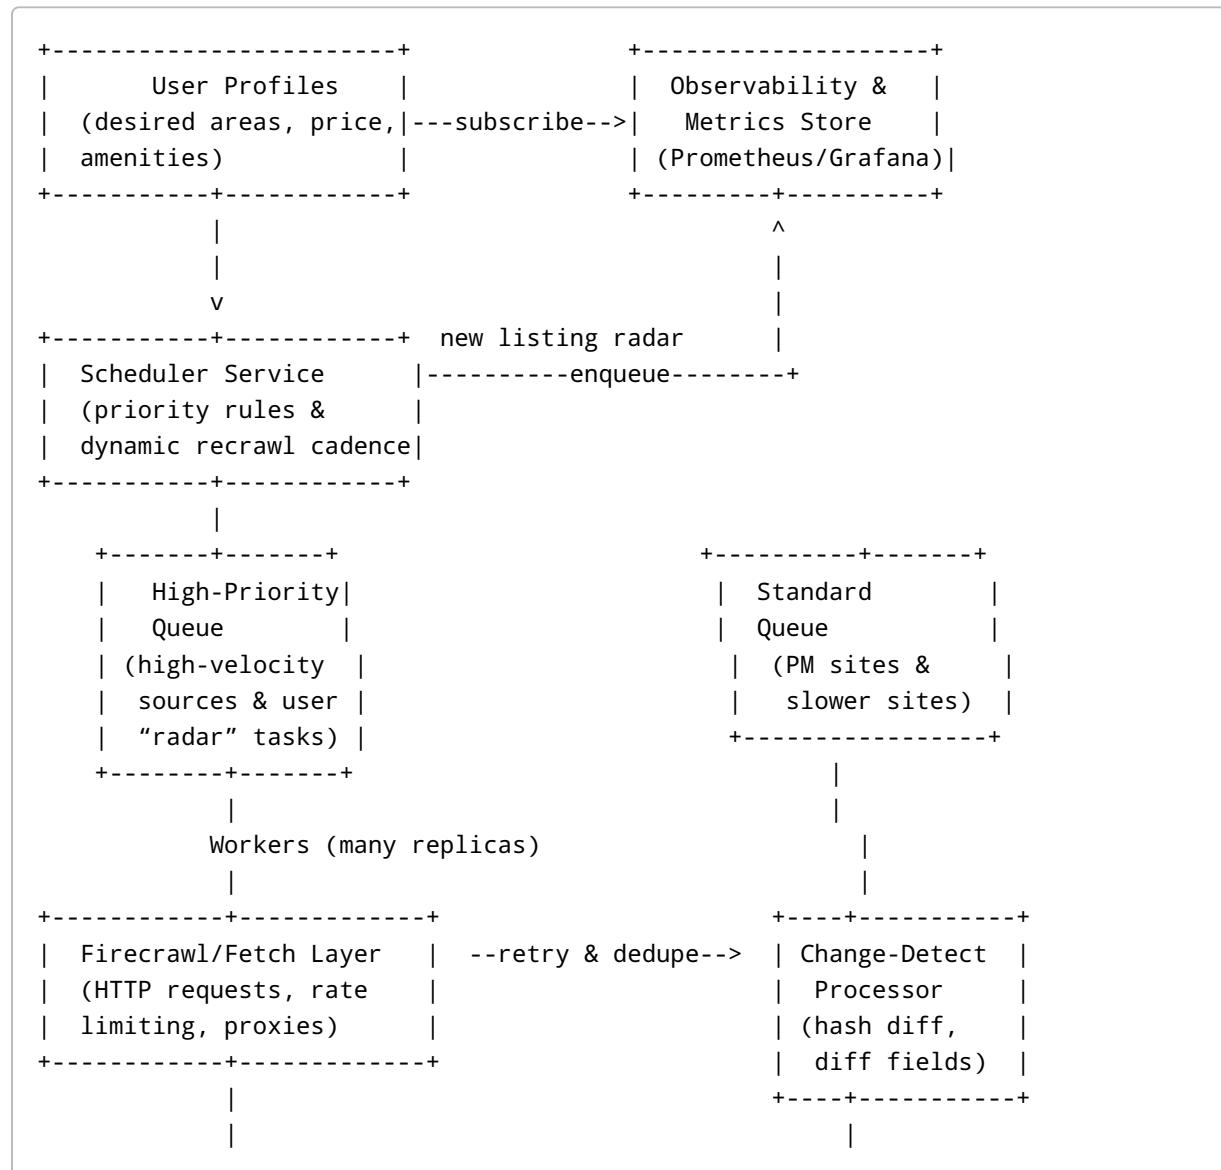


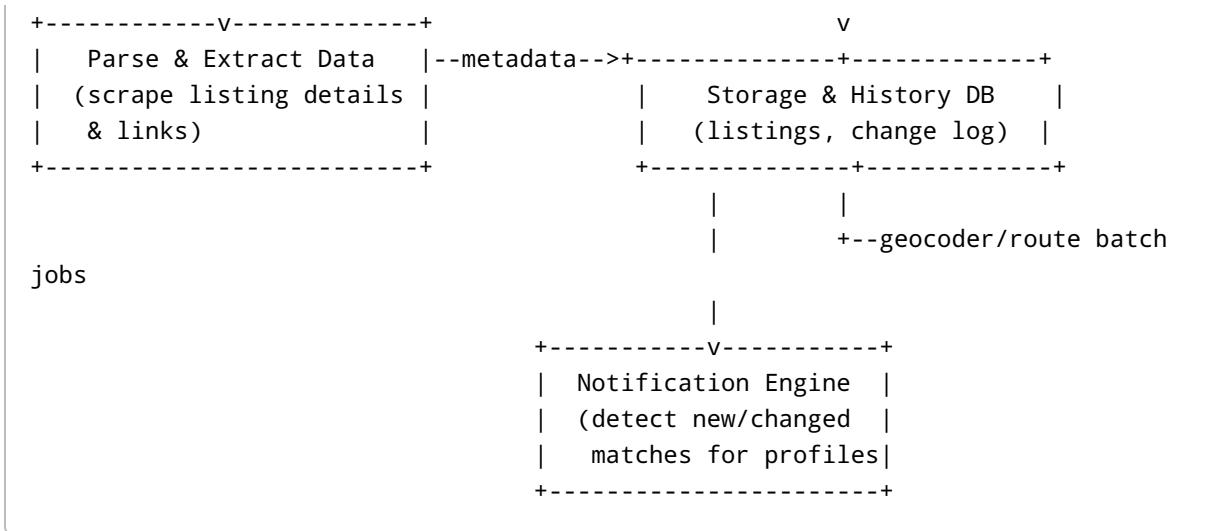


# Continuous Apartment-Listing Discovery System Design

This proposal describes a continuous discovery system for a personal apartment-hunting tool serving San Francisco. The user cares about finding new listings quickly, keeping the dataset fresh and broad, and robustly retrying failures. The design is not aimed at enterprise-grade security but emphasises correctness, freshness, and operational reliability.

## A) Architecture sketch





## Components

1. **Scheduler Service** – central brain that continuously decides what to crawl next. It maintains per-source metadata (last crawled time, historical change rate, failure counts) and uses dynamic recrawl policies.
  2. **Queues** – separate queues for high-velocity sources (Craigslist, Facebook Marketplace, etc.) and slower property-management sites. A special *radar* queue handles user-defined high-priority queries or neighbourhoods.
  3. **Workers** – stateless fetchers (async tasks or serverless functions) that pop URLs from queues, fetch pages using **Firecrawl** or a headless browser, enforce politeness, handle retries/exponential back-off, and push results to parsing.
  4. **Firecrawl/Fetch Layer** – interacts with Firecrawl API (or built-in fetcher) to request pages. It sets `changeTracking` and `markdown` formats to let Firecrawl compare current pages to previous versions; Firecrawl returns `changeStatus` (`new`, `same`, `changed`, `removed`) and timestamps
    - 1.
  5. **Parse & Extract** – interprets raw HTML / Firecrawl results into structured listing records (address, price, bedrooms, description, images). New URLs discovered on listing pages are normalised and deduped before enqueueing.
  6. **Change-Detection Processor** – compares the new record to the previous version. It uses field-level hashing and diffing. For each listing, compute hashes of price, text description, amenity fields and image URLs; when any hash changes, log a change event. Use tree-based hashing for structural changes to pages as described by Goel & Aggarwal: assign hash values to leaf nodes and tag values to non-leaf nodes in a DOM tree; compute bottom-up to identify modified nodes
    - 2.
  7. **Storage & History DB** – stores current listing snapshot plus a history table capturing each change event with timestamp and diff summary. Provide APIs for queries like “show price history for listing X.”
  8. **Notification Engine** – monitors change events and user profiles. When a new listing or a significant change matches a user’s criteria, it sends alerts (“Your top profile has 3 new matches since 9 AM”).
  9. **Observability & Metrics** – collects metrics (crawl coverage, freshness, error rates, queue sizes) and displays dashboards and alerts.

## B) Crawl scheduling rules

Scheduling decides what to crawl next. The **URL Frontier** should support multiple strategies:

1. **Priority scoring** – Each URL or source gets a score based on importance and likelihood of change.  
Use signals like:
2. **Historical change frequency** – pages that change often should be revisited sooner <sup>3</sup>. Track last-modified times or Firecrawl's `previousScrapeAt` and adapt.
3. **Source velocity** – classify sources as *high-velocity* (e.g., Craigslist, Zillow, Facebook groups) or *slow* (property-management sites). High-velocity sources receive shorter crawl intervals (minutes or hours). Slower sources might be recrawled daily or weekly.
4. **User interest** – if many users watch a neighbourhood or a price band, increase priority for listings in that segment (the **radar** queue). Provide “new listing radar” mode that polls specific sources/queries with high frequency during business hours.
5. **Listing age** – new pages discovered recently get higher scores for the first day, then gradually decrease as they become stale.
6. **Failure/back-off** – penalise sources with repeated errors (4xx/5xx). Use exponential back-off on error counts.
7. **Separate per-domain queues** – to respect politeness, maintain domain-level buckets so each domain is crawled independently with rate limits <sup>4</sup>. Limit concurrent requests per domain.
8. **Revisit policy** – compute a *next crawl time* per page based on its last change time and change frequency distribution. For example, if a listing’s price changed 4 times in the last week, schedule more frequent recrawls. For pages that haven’t changed in months, schedule recrawls weekly or monthly. Use time-based scheduling as recommended in incremental crawlers <sup>5</sup>.
9. **Deduplication & normalization** – before enqueueing a URL, normalise (remove # fragments, resolve relative links) and check a Bloom filter or fingerprint database to avoid duplicates <sup>6</sup>. Use SimHash or similar to detect near-duplicate pages <sup>7</sup>.
10. **Back-pressure & batch limits** – when queues grow, temporarily lower the crawling rate to control costs. Prioritise tasks with highest scores and drop extremely low-value tasks (e.g., old pages outside user’s price bands).

Pseudo-code for scoring:

```
score(url) = w_source * source_velocity_factor + w_change *  
estimated_change_prob \  
+ w_interest * user_interest_factor - w_age * page_age - w_fail *  
error_penalty
```

Jobs with higher scores are popped first. The scheduler continuously updates scores as new information arrives.

## C) Change detection design

1. **Hash-based field diffs** – For each listing, compute a hash of each field (price, availability, bedrooms, description text, images). Store the hashes and values. On subsequent crawls, recompute and compare. If hashes differ, record a change event with `old_value`, `new_value`, timestamp, and difference type (price-drop, description change, photo update, etc.).
2. **Tree-based change detection** – For pages where structural changes matter (e.g., a property page's layout), model the DOM as a tree and assign hash values to leaf nodes and tag values to non-leaf nodes; compute bottom-up to identify modified nodes as described by Goel & Aggarwal <sup>8</sup>. This approach is efficient for capturing both structural and content changes.
3. **Firecrawl changeTracking** – Use Firecrawl's built-in change-tracking mode to detect whether a page is `new`, `same`, `changed`, or `removed` <sup>9</sup>. When the status is `changed`, Firecrawl can return a `git-diff` or `json` diff showing modified text or structured fields. Use JSON mode to focus on price and availability fields for product-like pages <sup>10</sup>.
4. **Image change detection** – Compute perceptual hashes (pHash) of each photo. Compare new vs previous pHashes; if the Hamming distance exceeds a threshold, mark the photo as changed.
5. **Time series** – Record change events in a history table keyed by listing ID. Each entry stores the diff summary and timestamp. This allows analytics on price trajectories and helps predict future changes.

## D) Freshness metrics and alerts

To ensure the dataset stays current and that high-priority profiles surface new matches quickly, the system should collect metrics:

Metric	Purpose / Description
<b>Coverage</b>	Number/percentage of known sources crawled in the last $N$ hours (per domain, per neighbourhood). Identify coverage holes when certain areas or price ranges have not been updated recently.
<b>Freshness score</b>	Average age of listings (time since last crawled). Compute per-listing and overall distribution. Use as a target (e.g., 95 % of listings updated within 24 h).
<b>Change detection rate</b>	Ratio of crawls that resulted in a change event. High rates on a source indicate it is "hot" and may justify higher crawl frequency.
<b>Error budget</b>	Track HTTP errors, parsing failures, and timeout rates. Set an error budget (e.g., error rate <2 %). Alerts fire when error budgets are exceeded so engineers can investigate.
<b>Queue lag</b>	Time tasks spend in the queue before being processed. Long lag indicates insufficient worker capacity.
<b>User alert metrics</b>	Number of new/changed listings delivered to each user profile per day. Use this to provide notifications like "your top profile has 3 new matches since 9 AM".

Alerts and dashboards (e.g., using Grafana) can display these metrics and highlight anomalies. For example, if the *freshness score* for a high-velocity source exceeds 3 hours, an alert triggers to accelerate its crawl cadence.

## E) MVP (2-week) vs ultimate (2-3-month) roadmap

### Minimal viable product ( $\approx$ 2 weeks)

- **Scope** – Focus on a handful of high-velocity sources (e.g., Craigslist SF apartments, Zillow, a couple of property-management sites). Ignore dynamic rendering and anti-bot circumvention for now.
- **Tech stack** – Use a serverless or simple queue/worker architecture (e.g., cron-based scheduler, Firecrawl API) running on one machine. Store listings and change history in a relational DB (SQLite/PostgreSQL). Build a minimal web UI to show current listings and change logs.
- **Scheduling** – Hard-coded crawl intervals: high-velocity sources every 30 min, slower sources once daily. Basic deduplication using a Bloom filter. Use Firecrawl's `changeTracking` to detect new/changed/removed pages <sup>1</sup>.
- **Change detection** – Hash only key fields (price, status, description). Use Firecrawl JSON mode to track price and availability changes <sup>10</sup>. Log changes; send simple email notifications when a listing's price drops or a new listing matches a user's saved search.
- **Observability** – Collect simple metrics: number of crawled pages per source per day, error counts, time since last crawl. Use console logs or a lightweight dashboard.
- **Batch geocoding** – Geocode addresses nightly using a bulk geocoding API; store lat/long for mapping. Perform route distance calculations offline.

### Ultimate system ( $\approx$ 2-3 months)

- **Scalable architecture** – Deploy the Scheduler, Workers, and Storage as microservices on Kubernetes or serverless (AWS Lambda). Use distributed message queues (Kafka/SQS). Support thousands of URLs per minute.
- **Sophisticated scheduling** – Implement dynamic scoring based on historical change frequency, user interest signals, and error rates. Use machine-learning models to predict when a listing is likely to change and adjust crawl cadence accordingly. Implement per-domain rate-limiting buckets and polite crawling <sup>4</sup>.
- **Robust change detection** – Use tree-based hashing and Firecrawl's git-diff/JSON diff to capture small text edits, price updates, availability changes, and image updates. Use perceptual hashing for images. Maintain full time-series histories for each field.
- **Incremental crawling** – Instead of re-fetching entire pages, use HTTP `If-Modified-Since` or `ETag` headers to fetch only when content changes <sup>11</sup>. Only parse changed pages. Persist state and only re-scraper new entries since last run <sup>12</sup>.
- **Broad coverage** – Expand to dozens of property-management sites, listing aggregators, and community groups. Build custom adapters for sites requiring JavaScript rendering or anti-bot measures. Use proxies and rotate IPs to avoid bans.
- **Performance & cost optimization** – Cache responses from unchanged pages; dedupe early to avoid duplicate fetches <sup>6</sup>. Batch geocoding and route computations; precompute features used for search and ranking. Use incremental backups and compression for storage.

- **Observability** – Implement comprehensive dashboards with coverage, freshness, change detection rate, error budgets, and queue metrics. Use logs and traces to debug failures. Provide per-neighbourhood/price band coverage heatmaps to identify holes.
- **User experience** – Build a rich front-end with search filters, map view, price history charts, and alert preferences. Use asynchronous notifications (email, SMS, push) triggered by the Notification Engine when new matches appear.

## Practical cost/performance tactics

Despite an unlimited budget, good engineering practices help reduce latency and cost:

- **URL deduplication and normalization** – Use hashing (e.g., SimHash) to avoid crawling identical pages <sup>6</sup>.
- **Incremental crawling** – Use `If-Modified-Since` / `ETag` headers to fetch only updated content <sup>11</sup>. For API-based sources, use `updated_since` query parameters <sup>13</sup> to fetch only listings that changed since a certain time.
- **Partial parsing** – For sources with feed endpoints, request only metadata (e.g., JSON) rather than full HTML <sup>14</sup>.
- **Exponential back-off & retries** – Implement retry policies that back off on failure to reduce wasted calls <sup>15</sup>.
- **Batch geocoding & routing** – Instead of calling the geocoder for each address individually, accumulate addresses and call the geocoding API in bulk to take advantage of volume discounts.
- **Cache & CDN** – Cache images and static resources in a CDN to speed up repeated access. Use local caches to avoid re-downloading identical pages.

## Conclusion

This design provides a robust, always-on discovery system tailored for the user's apartment-hunting needs. It balances freshness and coverage by using dynamic scheduling and high-velocity radar queues, employs efficient change-detection techniques to track price and availability changes, and includes practical cost/performance optimizations. Observability and metrics ensure the system remains reliable, while the roadmap shows how to evolve from a simple MVP into a comprehensive platform over a few months.

---

1 9 10 Change tracking | Firecrawl

<https://docs.firecrawl.dev/features/change-tracking>

2 8 An Efficient Algorithm for Web Page Change Detection

<https://www.ijcaonline.org/archives/volume48/number10/7386-0173/>

3 What is a web crawler and how does it work? Full guide. | Parallel Web Systems | Web Search & Research APIs Built for AI Agents

<https://parallel.ai/articles/what-is-a-web-crawler>

4 6 7 15 Design a Web Crawler System Design: A Complete Guide

<https://www.systemdesignhandbook.com/guides/design-a-web-crawler-system-design/>

5 11 14 What strategies do incremental web crawlers use when processing web links? - Tencent Cloud

<https://www.tencentcloud.com/techpedia/115223>

12 13 How to perform Incremental Web scraping

<https://stabler.tech/blog/how-to-perform-incremental-web-scraping>